

CSCI 104

Hash Tables Intro

Mark Redekopp

David Kempe

Motivation

Suppose a company has a unique 3-digit ID for each of its 1000 employees.

- We want a data structure that, when given an employee ID, efficiently brings up that employee's record.

How should we implement this?


- An array gives $O(1)$ access time!

Alright, how do we obtain this runtime when the keys are no longer so nicely ordered or non-integers??

Maps/Dictionaries

Arrays

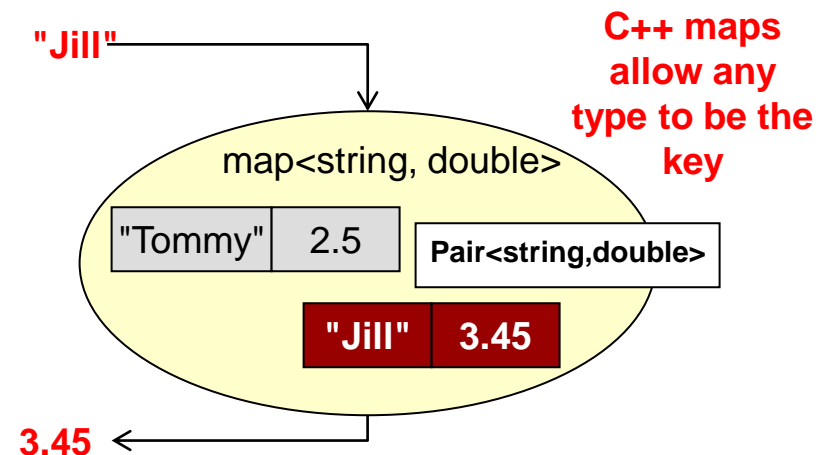
- An array maps integers to **values**
 - Given i , $\text{array}[i]$ returns the value in $O(1)$



Arrays associate an integer with some arbitrary type as the value (i.e. the key is always an integer)

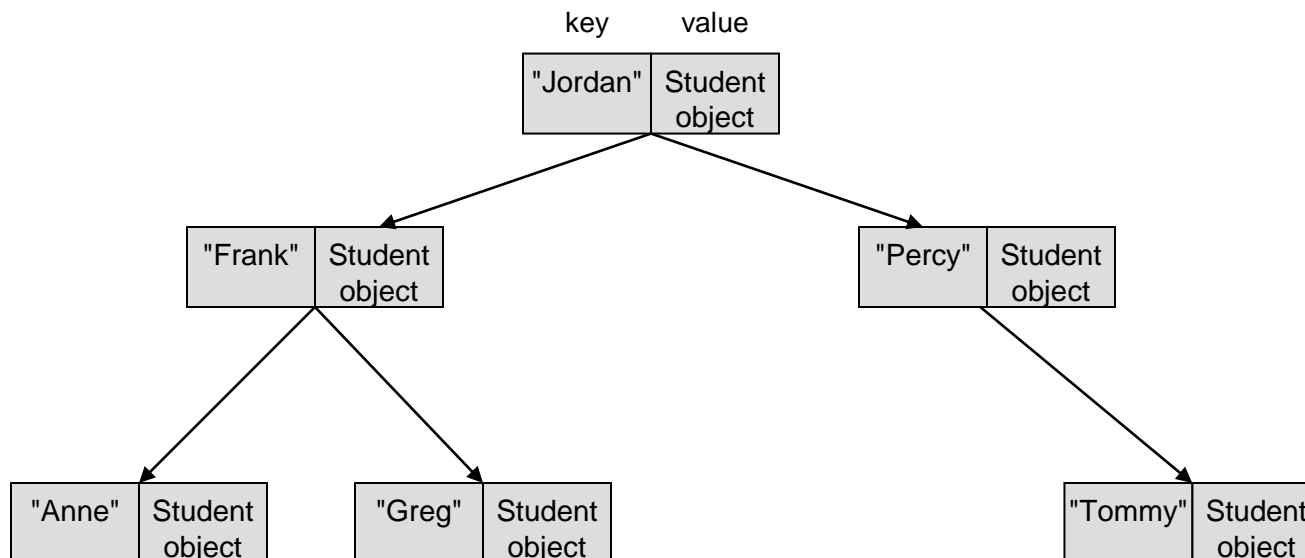
Maps/Dictionaries

- Dictionaries map keys to **values**
 - Given key, k , $\text{map}[k]$ returns the associated value
 - Key can be anything provided...
 - It has a ' $<$ ' operator defined for it (C++ map) or some other comparator functor (other languages require something similar)



Dictionary Implementation


- A dictionary/map can be implemented with a balanced BST
 - Insert, Find, Remove = $O(\text{_____})$
- Can we do better?
 - Hash tables (unordered maps) offer the promise of $O(\text{_____})$ access time



Hash Tables - Insert

- Can we use non-integer keys to index an array?
- Yes. Let us convert (i.e. "hash") the non-integer key to an integer
- To **insert** a key, we hash it and place the key (and value) at that index in the array
 - For now, make the unrealistic assumption that each unique key hashes to a unique integer
- The conversion function is known as a **hash function, $h(k)$**
- A hash table implements a **set/map** ADT
 - `insert(key)` / `insert(key,value)`
 - `remove(key)`
 - `lookup/find(key) => value`
- **Question to address:** What should we do if two keys ("Jill" and "Erin") hash to the same location (aka a **COLLISION**)?

`insert("Jill",3.7)`



A map implemented as a hash table
(key=name, value = GPA)

Hash table parameter definitions:

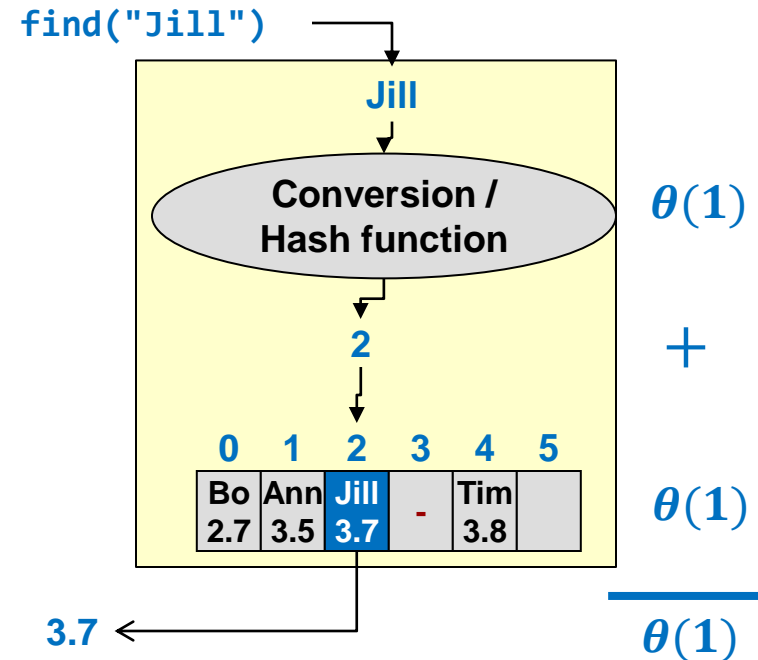
n = # of keys entered (=4 above)

m = tableSize (=6 above)

**$\alpha = \frac{n}{m}$ = Loading factor =
(4/6 above)**

Hash Tables - Find

- To **find** a key, we simply hash it again to find the index where it was inserted and access it in the array
- How might we hash a string to an integer?
 - Use ASCII codes for each character and **add, multiply, or shift/mix** them
 - We then can use simple a **modulo m** operation to convert the sum to a value between **0** to **$m-1$** where **m** is the table size
 - Note: All data in a computer is already bits (1s and 0s). Any object can be viewed as a long binary number and hashed



We could sum the ASCII values.

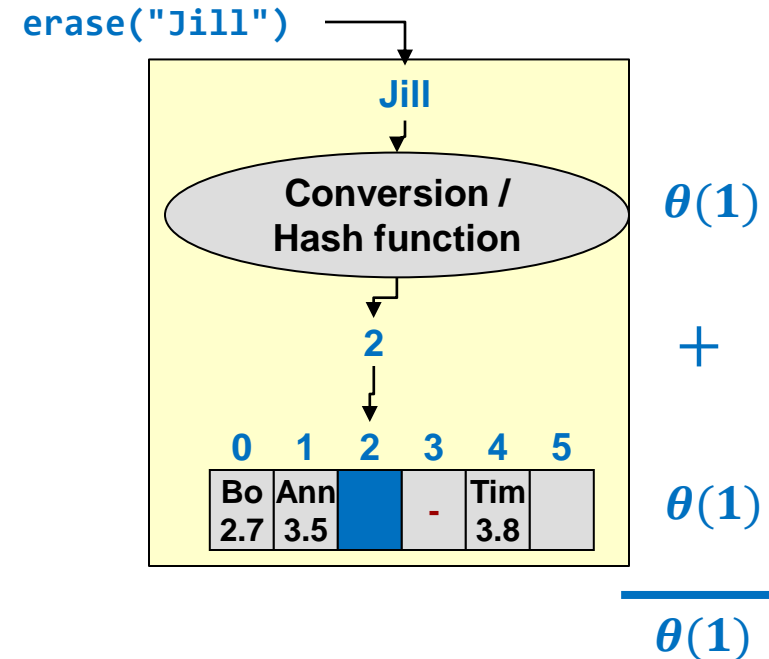
'h' = 104 'e' = 101 'l' = 108
 'l' = 108 'o' = 111

$h(\text{"hello"}) = 532 \% m$

Is this a good way to hash a string?

Hash Tables - Remove

- To **remove** a key, we simply hash the key and mark the location as "free" again
 - Could use a bool in the struct for each array entry (more later) to indicate it is free
- The **hash function, $h(k)$, should**
 - Be **fast/easy** to compute
 - $O(|k|)$ – where $|k|$ is the length of the key
 - But in terms of n [# of keys in the set/map] this runtime is constant since $|k| \ll n$ [e.g. $O(1)$]
 - Be **consistent** and output the same result any time it is given the same input
 - **Distribute** keys well
 - We'd like every unique key to map to a different index, but that turns out to be almost impossible.
 - We'll settle for a "**good**" hash function where the probability of a key mapping to any location x is $1/m$ (i.e. uniform)



Hash table parameter definitions:

n = # of keys entered

m = tableSize

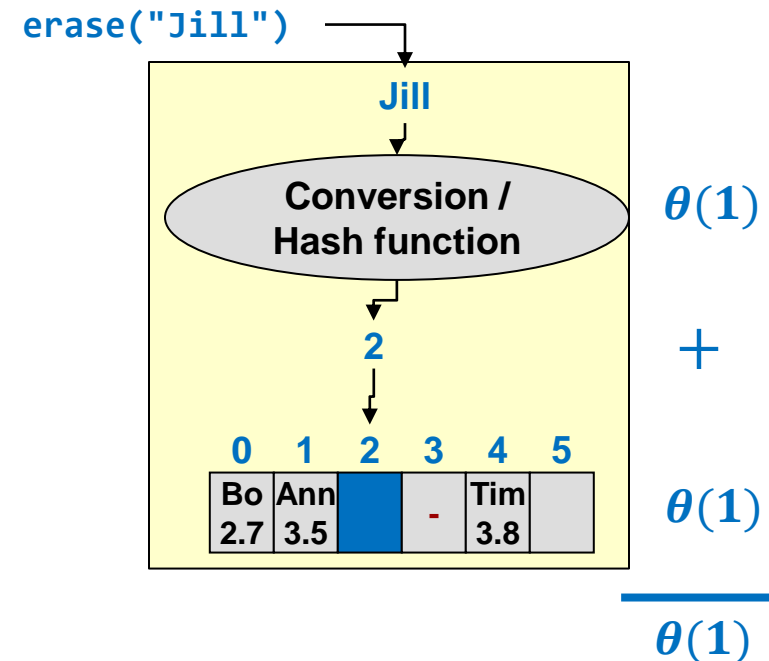
$\alpha = \frac{n}{m}$ = Loading factor

Possible Hash Functions

- Define n = # of keys stored, m = table size and suppose k is non-negative integer key
- Evaluate the following possible hash functions
 - $h(k) = 0$?
 - $h(k) = \text{rand}() \bmod m$?
 - $h(k) = k \bmod m$?
- Rules of thumb
 - The hash function should examine the entire search key (i.e. all bits/characters), not just a few digits or a portion of the key
 - When modulo hashing is used, the base should be prime

Hashing Efficiency

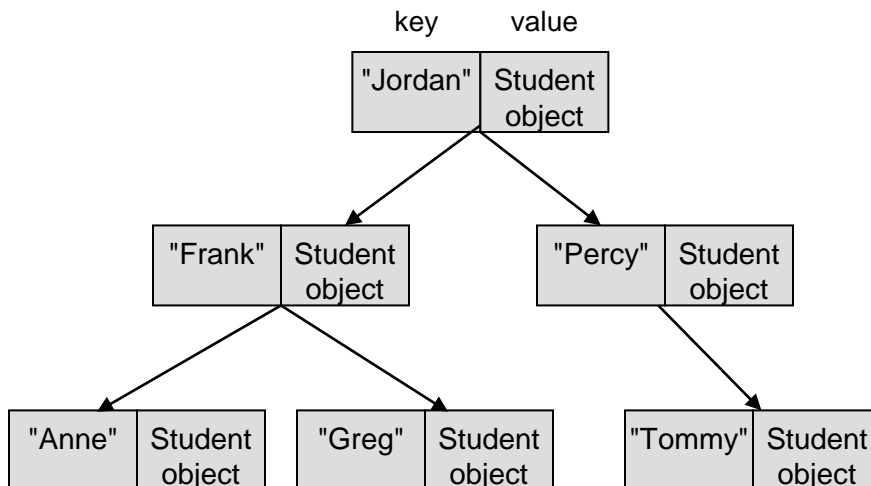
- If computing the hash function, $h(k)$, is $O(1)$ and the array access is $O(1)$,
- Then the runtime of the operations is $O(1)$
- What might prevent us from achieving this $O(1)$?
 - Collisions



Ordered vs. Unordered

Ordered Map/Set

- map/set
(implemented as balanced BST)
- $\log(n)$ runtime for insert/find/remove
- If we print each key via an in-order traversal of the tree, in what order will the keys be printed?



Unordered Map/Set

- unordered_map/unordered_set
(implemented as hash table)
- Each uses a hash table for $O(1)$ average runtime to insert, find, and remove
- New to C++11 and requires compilation with the `-std=c++11` option in g++
- Iteration will print the keys in an undefined order (unordered)
- Provides hash functions for basic types: int, string, etc. but for any other type you must provide your own hash function (like the operator< for BSTs)

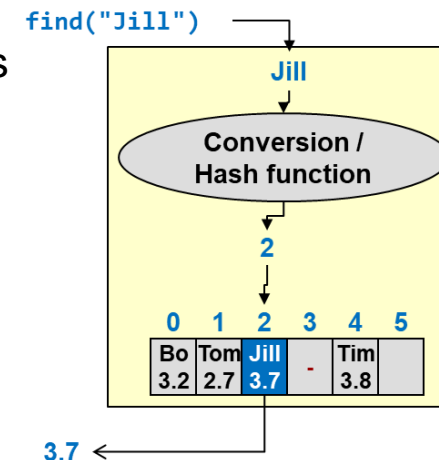
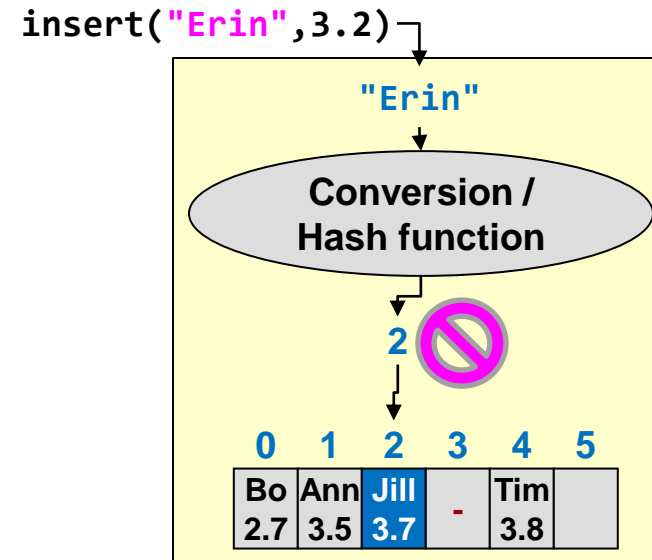


Table Size and Collisions

- Suppose we want to store USC student info using their 10-digit USC ID as the key
 - The set of all POSSIBLE keys, S , has size $|S| = 10^{10}$
 - But the number of keys we'd actually store, n , is likely much less (i.e. $n \ll |S|$)
- So how large should the table size (m) be?

_____ < _____ < _____

- But anything smaller than the size of all possible keys admits the chance of **COLLISION**
 - A collision is when two keys map to the same location [i.e. $h(k1) == h(k2)$]
 - The **probability** of this should be low
 - How we handle collisions is the major remaining question to answer
- You will see that table size (m) should usually be a **prime** number



COLLISION!!
 $h(\text{"Jill"}) = h(\text{"Erin"})$

Resolving Collisions


- Collisions occur when two keys, k_1 and k_2 , are not equal, but $h(k_1) = h(k_2)$.
- Collisions are inevitable if the number of entries, n , is greater than table size, m (**by pigeonhole principle**) and are likely even if $n < m$ (**by the birthday paradox...more in our probability unit**)
- Methods
 - **Closed Addressing** (e.g. buckets or **chaining**): Keys MUST live in the location they hash to (thus requiring multiple locations at each hash table index)
 - Methods: 1.) Buckets, 2.) Chaining
 - **Open Addressing (aka probing)**: Keys MAY NOT live in the location they hash to (only requiring a single 1D array as the hash table)
 - Methods: 1.) Linear Probing, 2.) Quadratic Probing, 3.) Double-hashing

Closed Addressing Methods

- Make each entry in the table a fixed-size ARRAY (bucket) or LINKED LIST (chain) of items/entries so all keys that hash to a location can reside at that index
 - Close Addressing** => A key **will reside in the location it hashes to** (it's just that there may be many keys (and values) stored at that location)
- Buckets**
 - How big should you make each array?
 - Too much wasted space
- Chaining**
 - Each entry is a linked list (or, potentially, vector)


Bucket 0

k,v				
0	Tim			...
1				...
2	Jill	Erin		...
3				...
4				...
				...
m-1	Bo			...




Open Addressing and Linear Probing

- With **open addressing**, we keep the hash table a 1D array (only one location per index) but when collisions occur we allow keys to reside in a location other than $h(k)$
 - **Open Addressing** => A key **may NOT reside in the location it hashes to** requiring extra searching in a process called **probing**
- For insertion: always start by checking location $h(k)$
 - If it is open, write the key (and value) there
 - Else "**probe**" for an empty location
- **Linear Probing (other techniques in a minute)**
 - Let **i** be number of **failed** checks to find a blank location (for insertion) or the key we are looking for (for find/remove)
 - $h(k, i) = (h(k) + i) \bmod m$
 - Example: If $h(k)$ occupied (i.e. collision) then check $h(k)+1, h(k)+2, h(k)+3, \dots$




Probing Impact on Find

- If $h(k)$ is occupied with another key, then probe
- **Insert:** probe until we find a blank location
- **Find/Remove:** probe until we...
 - Find the key we are looking for **..OR..**
 - _____ **..OR..**
 - _____





Probing Impact on Find

- If $h(k)$ is occupied with another key, then probe
- **Insert:** probe until we find a blank location
- **Find/Remove:** probe until we...
 - Find the key we are looking for **..OR..**
 - We reach a free location **..OR..**
 - We have looked in all possible locations (i.e. wrapped back to $h(k)$ or alternatively we've performed m probes)




Removal

- Many implementations exist but we will show one simple way for illustration
- Each location stores two bools
 - Valid**: a stored key exists in this location (or else is free)
 - Removed**: a key was erased at this location (so it is free for insertion, but probing must continue for find/remove)
- Progression:
 - Initially: **V=0, R=0** (Free/Never used),
 - On insert: **V=1, R=0**,
 - On erasure: **V=0, R=1** (can return to V=1, R=0 on insert)
- For performance, we can periodically rebuild/rehash the hash table after some number of erasures to effectively return locations to **free/never used**




Linear Probing & Primary Clustering

- Suppose a hash table ($m=10$) with integer keys and $h(k) = k \% m$
- Insert: 11, 21, 2, 31, 3
 - Notice, that the collisions of 11, 21, and 31 cause collisions for 2 and 3 which then may cause collisions for other nearby hash locations
- This is known as **primary clustering** (a few collisions to one location and the resulting probing cause collisions for other keys that would not have collided)





Quadratic Probing

- If certain data patterns lead to many collisions, linear probing leads to clusters of occupied areas in the table called **primary clustering**
- **Quadratic probing** tends to spread out data across the table by taking larger and larger steps until it finds an empty location
- **Quadratic Probing**
 - (Again, let i be number of **failed** probes)
 - $h(k, i) = (h(k) + i^2) \bmod m$
 - If $h(k)$ occupied, then check $h(k) + 1^2$, $h(k) + 2^2$, $h(k) + 3^2$, ...



Linear vs. Quadratic Probing

- If certain data patterns lead to many collisions, linear probing leads to clusters of occupied areas in the table called ***primary clustering***
- How would quadratic probing help fight primary clustering?
 - Quadratic probing tends to spread out data across the table by taking larger and larger steps until it finds an empty location



Quadratic Probing Practice

- Use the hash function $h(k)=k\%9$ to find the contents of a hash table ($m=9$) after inserting keys 36, 27, 18, 9, 0 using quadratic probing
- If your **loading factor** rises above 0.5, bad things can happen!

0	1	2	3	4	5	6	7	8


- Use the hash function $h(k)=k\%7$ to find the contents of a hash table ($m=10$) after inserting keys 14, 8, 21, 2, 7 using quadratic probing

0	1	2	3	4	5	6

- Quadratic probing only works well for prime table sizes, and keeping the load factor < 0.5

Double Hashing

- Note: In linear and quadratic probing, if two keys hash to the same place ($h_1(k_1) == h_1(k_2)$) we will probe the **same** sequence
- Could we probe a **different** sequence even if two keys have collided?
 - Let's use ANOTHER hash function, $h_2(k)$ to choose the step size of our probing sequence
- Double Hashing**
 - (Again, let i be number of **failed** probes)
 - Pick a second hash function $h_2(k)$ in addition to the primary hash function, $h_1(k)$
 - $h(k, i) = [h_1(k) + i * h_2(k)] \bmod m$



Sequence:

- Start at $h_1(k)$,
- If needed, probe $h_1(k) + h_2(k)$
- If needed, probe $h_1(k) + 2 * h_2(k)$
- If needed, probe $h_1(k) + 3 * h_2(k)$

Double Hashing

- Assume
 - $m=13$,
 - $h_1(k) = k \% 13$
 - $h_2(k) = 5 - (k \% 5)$
- What sequence would I probe if $k = 31$
 - $h_1(31) = 5$
 - $h_2(31) = 5 - (31 \% 5) = 4$ (which is the step size)
 - $5 + 0*4 = 5 \% 13 = 5$
 - $5 + 1*4 = 9 \% 13 = 9$
 - $5 + 2*4 = 13 \% 13 = 0$
 - $5 + 3*4 = 17 \% 13 = 4$
 - And then onto 8, 12, 3, 7, 11, 2, 6, 10, 1
 - Notice we visited each index in the table. Why? **A prime table size!**

Rehashing

- For probing (open-addressing), as α approaches 1 the expected number of probes/comparisons will get very large
 - Capped at the tableSize, m (i.e. $O(m)$)
- Similar to resizing a vector, we can allocate a larger prime size table/array
 - Must **rehash** items to location in new table size and **cannot just copy items to corresponding location in the new array**
 - Example: $h(k) = k \% 7 \neq h(k) = k \% 11$ (e.g. $k=9$)
 - For quadratic probing if table size m is prime, then first $m/2$ probes will go to unique locations
- General guideline for probing: keep $\alpha < \underline{\hspace{1cm}}$

0	1	2	3	4	5	6		0	1	2	3	4	5	6	7	8	9	10
	1	9	38	18														



$$h(k) = k \% 7$$

$$h(k) = k \% 11$$

Rehashing

- For probing (open-addressing), as α approaches 1 the expected number of probes/comparisons will get very large
 - Capped at the tableSize, m (i.e. $O(m)$)
- Similar to resizing a vector, we can allocate a larger prime size table/array
 - Must **rehash** items to location in new table size and **cannot just copy items to corresponding location in the new array**
 - Example: $h(k) = k \% 7 \neq h(k) = k \% 11$ (e.g. $k=9$)
 - For quadratic probing if table size m is prime, then first $m/2$ probes will go to unique locations
- General guideline for probing: keep $\alpha < 0.5$**



0	1	2	3	4	5	6		0	1	2	3	4	5	6	7	8	9	10
	1	9	38	18					1				38		18		9	

$$h(k) = k \% 7$$

$$h(k) = k \% 11$$

Probing Technique Summary

- If $h(k)$ is occupied with another key, then probe
- Let i be number of **failed** probes
- **Linear Probing**
 - $h(k, i) = (h(k) + i) \bmod m$
- **Quadratic Probing**
 - $h(k, i) = (h(k) + i^2) \bmod m$
 - If $h(k)$ occupied, then check $h(k) + 1^2$, $h(k) + 2^2$, $h(k) + 3^2$, ...
- **Double Hashing**
 - Pick a second hash function $h_2(k)$ in addition to the primary hash function, $h_1(k)$
 - $h(k, i) = [h_1(k) + i * h_2(k)] \bmod m$



Hash Function Goals

- A "perfect hash function" should map each of the n keys to a unique location in the table
 - Recall that we will size our table to be larger than the expected number of keys...i.e. $n < m$
 - Perfect hash functions are not practically attainable
- A "good" hash function
 - Is easy and fast to compute
 - Scatters data uniformly throughout the hash table
 - $P(h(k) = x) = 1/m$ (i.e. **pseudorandom**)

Hashing Efficiency

- Loading factor, α , defined as:
 - $\alpha = n / m$ (Really it is just the fraction of locations currently occupied)
 - n =number of items in the table, m =tableSize
- For open addressing, $\alpha \leq 1$
 - Good rule of thumb: resize and rehash after $\alpha > 0.5$
- For closed addressing (chaining), α , can be greater than 1
 - This is because $n > m$
 - What is the average length of a chain in the table (e.g. 10 total items in a hash table with table size of 5)?
 - Need to keep α constant (usually $\alpha \leq 1$)

Hashing Efficiency

- Loading factor, α , defined as:
 - $\alpha = n / m$ (Really it is just the fraction of locations currently occupied)
 - n =number of items in the table, m =tableSize
- For open addressing, $\alpha \leq 1$
 - Good rule of thumb: resize and rehash after $\alpha > 0.5$
- For closed addressing (chaining), α , can be greater than 1
 - This is because $n > m$
 - What is the average length of a chain in the table (e.g. 10 total items in a hash table with table size of 5)?
 - Average length of chain will be $\alpha = n / m$
 - Need to keep α constant (usually $\alpha \leq 1$)

Hash Tables are Awesome!

Hash tables provide a very lucrative potential runtime. However, they are **probabilistic**.

- There was a similar problem with Splay Trees: they had a good **average** runtime, but a poor **worst**-case runtime.

As of this moment, we do not have the necessary mathematical framework to analyze either of these structures.

- We're going to start remedying that... now.